

CARD 1 OF 5

TIME: 30 MINS

TARGET: 10:00 AM

TASK 1: SOURCE INGESTION & LICENSE VERIFICATION**OBJECTIVES & LAB TASKS:**

- Acquire official guideline PDFs: WHO 2021 Hypertension & NICE NG238 Lipid Modification.
- Confirm source licensing (Open Access CC BY-NC-SA 3.0 IGO & Crown Copyright OGL).
- Verify file integrity, page counts, and absence of corrupted PDF text streams.

HARD STOP GUARDRAIL:

Stop source vetting by 10:00! Do NOT spend time reading guidelines. Move to parsing immediately.

DELIVERABLE:

data/raw/ directory populated with validated PDFs and confirmed licensing.

CARD 2 OF 5

TIME: 45 MINS

TARGET: 10:45 AM

TASK 2: EXTRACTION & MEDICAL CLEANING**OBJECTIVES & LAB TASKS:**

- Extract document text streams with page tracking using PyMuPDF.
- Extract structured clinical tables using pdfplumber.
- Strip repetitive running headers/footers while preserving medical terms.
- Preserve comparison operators ($\geq$, $\leq$, $<$, $>$) and dosage units (mmHg, mmol/L, mg).

HARD STOP GUARDRAIL:

Cleaning must be purely structural. Never modify clinical sentences or thresholds.

DELIVERABLE:

Cleaned page text stream with exact source PDF page boundaries intact.

CARD 3 OF 5

TIME: 60 MINS

TARGET: 11:45 AM

TASK 3: GUIDELINE STRUCTURE & RECOMMENDATION**OBJECTIVES & LAB TASKS:**

- Detect major section headings, subheadings, and rationale blocks.
- Parse discrete recommendation IDs (e.g. 1.1.4, 1.6.7, WHO Rec 1-8).
- Extract recommendation publication and amendment dates (e.g. [2014, amended May 2023]).
- Map cross-references, technology appraisals (TAs), and clinical algorithms.

HARD STOP GUARDRAIL:

Ensure 100% ID uniqueness across all chunks. Map orphan recommendations to parent sections.

DELIVERABLE:

Parsed document hierarchy tree with classified recommendation blocks.

CARD 4 OF 5

TIME: 45 MINS

TARGET: 12:30 PM

TASK 4: SEMANTIC CHUNKING & CLINICAL ENRICHMENT**OBJECTIVES & LAB TASKS:**

- Produce standalone recommendation, rationale, table, and algorithm chunks.
- Enforce token length bounds ($50 \leq \text{tokens} \leq 800$) with tiktoken cl100k_base.
- Extract clinical metadata: populations, risk tools (QRISK3), lipid targets, drug doses.
- Assign clinical priority weights (Priority 1: Active Recommendations).

HARD STOP GUARDRAIL:

No orphan sentences or micro-chunks (< 20 tokens). Keep each chunk fully self-contained.

DELIVERABLE:

Enriched semantic chunk objects ready for Validation Gates V1-V6.

CARD 5 OF 5

TIME: 30 MINS

TASK 5: VALIDATION GATES V1-V6 & TRAINER SIGN-OFF

TARGET: 13:00 (CUTOFF: 13:30)

OBJECTIVES & LAB TASKS:

- Run Quality Gates V1-V6 (Unique IDs, Completeness, Bounds, Provenance, Priority, Idempotency).
- Emit final data/processed/*_chunks.json, *.jsonl, and QA preview markdown.
- Execute full pytest test suite (119/119 unit tests green).
- Present Readiness Checklist to Trainer for sign-off before 13:30 lunch cutoff!

HARD STOP GUARDRAIL:

13:30 LUNCH HARD CUTOFF. Trainer sign-off is required before afternoon index-building session.

DELIVERABLE:

Green test suite + validated chunks + signed Trainer Readiness Checklist.

IMPORTANT PACING REMINDER FOR DAY 1

1. Keep Table Focus: Display the current module timer card prominently at your workspace table.
2. Avoid Scope Creep: Do not spend time researching clinical guidelines beyond the scope of ingestion.
3. Deterministic Ingestion: Use rule-based parsing for 100% reproducibility and clinical safety.
4. Validation Gates V1-V6: All 6 quality gates must pass before requesting trainer sign-off.
5. 13:30 Lunch Cutoff: Ensure your pipeline produces validated JSON/JSONL outputs before lunch.